

Rishabh Verma

973-975-5796 | rishworkn@gmail.com | Newark, NJ | US Citizen
[linkedin.com/in/rishabh-verma01](https://www.linkedin.com/in/rishabh-verma01) | github.com/rishiaman2 | [Portfolio](#)

EDUCATION

New Jersey Institute of Technology – Ying Wu College of Computing

Newark, NJ

Master of Science in Data Science

Jan 2026 – May 2027

- **Relevant Coursework:** DS644 Big Data Systems, DS675 Machine Learning, DS637 Python and Mathematics for Data Science, DS680 NLP, DS667 Deep Learning, CS683 Software Project Management

New Jersey Institute of Technology – Ying Wu College of Computing

Newark, NJ

BS in Data Science (BS/MS track), Minor in Business | GPA: 3.3/4.0

Sep 2021 – Dec 2025

- **Relevant Coursework:** CS288 Intensive Programming in Linux, CS331 Database System Design and Management, CS482 Data Mining, CS370 Intro to Artificial Intelligence

EXPERIENCE

Forward Deployed Engineer Intern — Cloud Business Division (MSP)

Jun 2026 – Aug 2026

Samsung SDS America (SDSA)

Ridgefield Park, NJ

- Applied AI/ML expertise to equipment monitoring, failure prediction, and fault diagnostics; decoded system logic and outputs and connected technical components into a coherent end-to-end solution
- Worked with AWS-based ETL pipelines, Kafka-based event streaming, and data lake / object storage (S3-compatible) components to move equipment-sensor data into the monitoring and prediction pipeline
- Explored retrieval-augmented generation (RAG) and Model Context Protocol (MCP)-based integrations while driving prompt engineering and rapid prototyping on AWS, designing, testing, and iterating on prompts to turn concepts into working prototypes
- Promoted explainable and responsible AI practices — including PII-aware data handling and data-lineage/traceability discussions — so predictions stayed understandable, trustworthy, secure, and governed rather than opaque automated decisions
- Evaluated Snowflake and Databricks as candidate platforms for scalable data engineering and model experimentation, and supported FinOps-style reviews of AWS consumption using Apptio to help inform architecture and investment decisions

Software & Product Manager

Jan 2022 – Dec 2024

Invasion Media

Indore, India

- Built 3 dynamic versions of the company website's front-end and back-end feature integration
- Designed NDA-protected backend software, managed a development team of 5, and led debugging and deployment
- Introduced internal systems to track productivity and eliminate redundancy across teams by 40%

Finance Research Assistant

May 2024 – Feb 2025

Martin Tuchman School of Management, NJIT

Newark, NJ

- Applied the BEKK-GARCH model to analyze data from the Bloomberg Index and wrote the paper in LaTeX
- Co-authored (second author, 60% contribution) a decentralized finance research paper under Dr. Wang: *Do Global Uncertainty Factors Matter More to Cryptocurrency?* and presented the paper at SWFA 2025 in Texas

PROJECTS

ECG Data Processing with Hadoop MapReduce on AWS | Java, HDFS, YARN, EC2

Spring 2026

- Deployed a fully distributed 4-node Hadoop cluster (1 NameNode, 3 DataNodes) on AWS EC2 Ubuntu instances, installing and configuring Java 8 and Hadoop 3.x across all nodes
- Configured HDFS, YARN, and passwordless SSH (core-site, hdfs-site, mapred-site, yarn-site) to enable fully distributed processing and inter-node communication
- Diagnosed cluster bottlenecks, documented troubleshooting steps, and identified performance optimization strategies

All-in-One (AIO) Bot – Shade Softwares | Automation Platform

Jan 2022 – May 2022

- Built a scalable automation platform enabling 300 users to execute multiple concurrent web requests while mitigating detection and throttling risks
- Led end-to-end product lifecycle, including architecture design, pricing strategy, customer onboarding, and technical support

TECHNICAL SKILLS

Languages & Scripting: Python, Java, SQL, Bash, C++

Cloud & Architecture: AWS (EC2, S3), Cloud Architecture, System Design, Docker, Linux, Git

Distributed Systems & Data: Hadoop, HDFS, YARN, MapReduce, ETL, Kafka, S3/MinIO, Data Lakes, Snowflake, Databricks

AI & Governance: RAG, MCP, Explainable AI, Responsible AI, PII Governance, Data Lineage

Engineering & Operations: Application Development, Debugging, Deployment, Performance Optimization, Operational Reliability, Root-Cause Troubleshooting, MS Project, Apptio (FinOps)